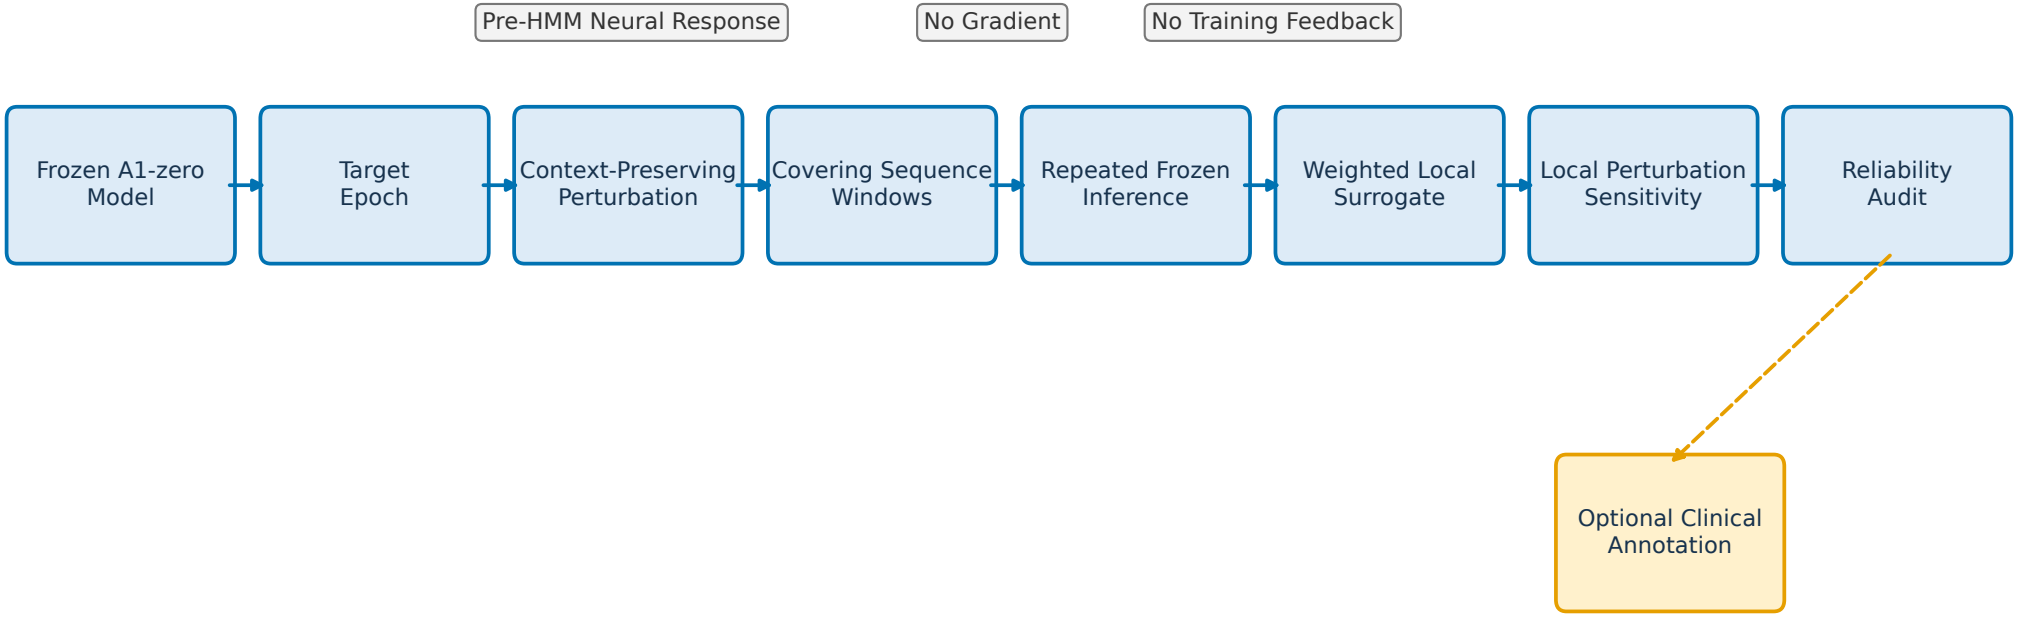
Pre-HMM Neural Response
No Gradient
No Training Feedback
Frozen A1-zero Model
Target Epoch
Context-Preserving Perturbation
Covering Sequence Windows
Repeated Frozen Inference
Weighted Local Surrogate
Local Perturbation Sensitivity
Reliability Audit
Optional Clinical Annotation
Query-only post-hoc path • normal inference never reads XAI